

CAF — Inputs Pipeline

Aggregated export — 2026-06-16. Source markdown in repo is unchanged.

Bundle id: 08-caf-inputs-pipeline

Included: docs/CAF_INPUTS_PIPELINE_ROADMAP.md

CAF inputs → signal pack → RTP roadmap

This backlog tracks **CAF Core** (APIs + Admin). The **Review** app stays focused on human review, publishing, and learning — not inputs processing operators.

Product split

Surface	Role
CAF Core Admin (/admin/inputs, /admin/processing; legacy /admin/inputs-processing → redirect)	Upload/history; pre-LLM + broad + top-performer (image, carousel, video) insights; profile, build signal pack, audit, RTP, QC flow profiles, insights packs.
Review → Pipeline	Light touch only: upload evidence XLSX into Core, browse imports, inspect signal-pack ideas for editorial context next to review work. No processing controls here.

Done (Core + Admin)

- **027** — inputs_evidence_imports / inputs_evidence_rows (XLSX ingest, dedupe keys).
- **063** — inputs_source_rows , inputs_scraper_config , inputs_scraper_runs — project source registry + Apify scraper runs into the same evidence import shape as XLSX.
- **028** — Row ratings + inputs_processing_profiles (criteria, models, caps, min score).
- **029** — Import/row health fields, selection snapshot on import, signal_packs.source_inputs_import_id , runs.plan_summary_json , insights_packs , qc_flow_profiles ; health + selection in build path.
- **Evidence upload API** — POST /v1/inputs-evidence/upload , list/detail/rows.

- **Processing API** — GET/PUT `.../profile` , import stats (`recompute_health`), POST `.../build-signal-pack` , audit, insights packs list, RTP summary, QC flow profiles CRUD.
- **Rating + synthesis** — OpenAI batch scoring → persisted rating columns → synthesis to `overall_candidates_json` (planner contract aligned with XLSX packs).
- **Admin UI — Inputs vs Processing** (sidebar); Processing segments: Evidence, broad insights per platform, top performers (image / carousel / video), profile & audit.
- **Inputs Admin** — tabs: **Imports** (XLSX upload unchanged), **Sources** (sidebar sheet picker + sync source tabs from workbook), **Scrapers** (Apify + HTML, run history → evidence import).
- **030-032** — `inputs_evidence_row_insights : tiers broad_llm , top_performer_deep , top_performer_video , top_performer_carousel` (migrations **030, 031, 032**). These tiers feed `visual_guidelines_pack_v1` used by **top-performer mimic** jobs — see [MIMIC FLOWS COMPLETE GUIDE.md](#).

Stage 2 (evidence insights) — complete

Row-level LLM passes are in place: text broad, single-image vision, multi-slide carousel vision, sampled-frame + transcript video. Admin can run each tier and list results by platform (`evidence_kind` on insights API).

Stage 3 (before signal pack) — ideas

Ideas are a structured list (many items) stored on or beside the signal pack. When a run creates **candidates**, the system mixes **project configuration constraints** with **selected ideas**: either an LLM chooses which ideas fit the run context, or a **human** picks them in admin/review. Implementation TBD (schema + UI + prompt contract).

Done (Review, minimal)

- Pipeline: **upload + list imports + inspect evidence rows; list + inspect signal-pack ideas** — proxies only for evidence/signal-pack reads (and upload). Processing endpoints are **not** mirrored in Review.

Next — richer scoring & planner audit (Core)

- **Overall idea scoring in Core** — Explainable breakdown beyond LLM pre-score; eligibility gates (brand, risk).
- **Persist planner selection** — Snapshot of `decideGenerationPlan` input/output per run (unify with existing decision traces).
- **Candidate row linkage** — Optional `caf_core.candidates` or exports from `generation_payload` for audit.

Next — summaries & RTP depth (Core)

- **HTML / platform summaries** — Optional pass: fold `HTML_Findings_Summary` + registry rows into `html_summary_json` / `reddit_summary_json` for richer planner context (non-breaking if null).
- **RTP by edit category** — Structured edit tracking (script / timing / template / assets) + dashboards vs carousel/video targets (beyond current strict RTP buckets).

Ops

- Run migrations: `npm run migrate` (or rely on `CAF_RUN_MIGRATIONS_ON_START`).